

```
# !pip install -q -U \
# transformers \
# datasets \
# accelerate \
# peft \
# trl \
# bitsandbytes \
# sentencepiece \
# safetensors \
# evaluate \
# scikit-learn \
# sacrebleu \
# rouge-score \
# bert-score \
# sentence-transformers
```

Preparing metadata (setup.py) ... done

```
===== 10.5/10.5 MB 72.4 MB/s eta 0:00:00
===== 529.0/529.0 kB 34.9 MB/s eta 0:00:00
===== 721.6/721.6 kB 41.6 MB/s eta 0:00:00
===== 60.7/60.7 MB 18.4 MB/s eta 0:00:00
===== 84.1/84.1 kB 5.5 MB/s eta 0:00:00
===== 8.9/8.9 MB 72.0 MB/s eta 0:00:00
===== 100.8/100.8 kB 5.3 MB/s eta 0:00:00
===== 61.1/61.1 kB 3.8 MB/s eta 0:00:00
===== 48.9/48.9 MB 13.9 MB/s eta 0:00:00
```

Building wheel for rouge-score (setup.py) ... done

```
# pip install cuda-toolkit==12.6.0
```

Collecting cuda-toolkit==12.6.0

Downloading cuda_toolkit-12.6.0-py2.py3-none-any.whl.metadata (8.3 kB)

Downloading cuda_toolkit-12.6.0-py2.py3-none-any.whl (2.2 kB)

Installing collected packages: cuda-toolkit

Attempting uninstall: cuda-toolkit

Found existing installation: cuda-toolkit 12.8.1

Uninstalling cuda-toolkit-12.8.1:

Successfully uninstalled cuda-toolkit-12.8.1

Successfully installed cuda-toolkit-12.6.0

```
# pip install torch==2.5.1+cu121 torchvision==0.20.1+cu121 torchaudio==2.5.1+cu121 -
```

```

    Successfully uninstalled sympy-1.14.0
Attempting uninstall: nvidia-nvtx-cu12
  Found existing installation: nvidia-nvtx-cu12 12.8.90
  Uninstalling nvidia-nvtx-cu12-12.8.90:
    Successfully uninstalled nvidia-nvtx-cu12-12.8.90
Attempting uninstall: nvidia-nccl-cu12
  Found existing installation: nvidia-nccl-cu12 2.27.5
  Uninstalling nvidia-nccl-cu12-2.27.5:
    Successfully uninstalled nvidia-nccl-cu12-2.27.5
Attempting uninstall: nvidia-cusparse-cu12
  Found existing installation: nvidia-cusparse-cu12 12.5.8.93
  Uninstalling nvidia-cusparse-cu12-12.5.8.93:
    Successfully uninstalled nvidia-cusparse-cu12-12.5.8.93
Attempting uninstall: nvidia-curand-cu12
  Found existing installation: nvidia-curand-cu12 10.3.9.90
  Uninstalling nvidia-curand-cu12-10.3.9.90:
    Successfully uninstalled nvidia-curand-cu12-10.3.9.90
Attempting uninstall: nvidia-cufft-cu12
  Found existing installation: nvidia-cufft-cu12 11.3.3.83
  Uninstalling nvidia-cufft-cu12-11.3.3.83:
    Successfully uninstalled nvidia-cufft-cu12-11.3.3.83
Attempting uninstall: nvidia-cuda-runtime-cu12
  Found existing installation: nvidia-cuda-runtime-cu12 12.8.90
  Uninstalling nvidia-cuda-runtime-cu12-12.8.90:
    Successfully uninstalled nvidia-cuda-runtime-cu12-12.8.90
Attempting uninstall: nvidia-cuda-nvrtc-cu12
  Found existing installation: nvidia-cuda-nvrtc-cu12 12.8.93
  Uninstalling nvidia-cuda-nvrtc-cu12-12.8.93:
    Successfully uninstalled nvidia-cuda-nvrtc-cu12-12.8.93
Attempting uninstall: nvidia-cuda-cupti-cu12
  Found existing installation: nvidia-cuda-cupti-cu12 12.8.90
  Uninstalling nvidia-cuda-cupti-cu12-12.8.90:
    Successfully uninstalled nvidia-cuda-cupti-cu12-12.8.90
Attempting uninstall: nvidia-cublas-cu12
  Found existing installation: nvidia-cublas-cu12 12.8.4.1
  Uninstalling nvidia-cublas-cu12-12.8.4.1:
    Successfully uninstalled nvidia-cublas-cu12-12.8.4.1
Attempting uninstall: nvidia-cusolver-cu12
  Found existing installation: nvidia-cusolver-cu12 11.7.3.90
  Uninstalling nvidia-cusolver-cu12-11.7.3.90:
    Successfully uninstalled nvidia-cusolver-cu12-11.7.3.90

```

```

from google.colab import drive
drive.mount("/content/drive")

```

Mounted at /content/drive

```

from pathlib import Path
import os
import sys
import json
import time
import gc

import torch
import pandas as pd

from datasets import load_dataset
from transformers import AutoTokenizer, AutoModelForCausalLM, BitsAndBytesConfig
from peft import PeftModel

```

```
from google.colab import userdata
```

```
os.environ["HF_TOKEN"] = userdata.get("HF_TOKEN")
print("HF_TOKEN cargado:", bool(os.environ.get("HF_TOKEN")))
```

```
HF_TOKEN cargado: True
```

```
DRIVE_ROOT = Path("/content/drive/MyDrive/TT2_colab")
```

```
DATA_DIR = DRIVE_ROOT / "data"
```

```
SFT_DIR = DATA_DIR / "sft_ready"
```

```
OUTPUT_DIR = DRIVE_ROOT / "outputs" / "lora_runs"
```

```
TEST_JSONL = SFT_DIR / "feina_repr30_test_sft.jsonl"
```

```
print("TEST_JSONL:", TEST_JSONL)
```

```
print("OUTPUT_DIR:", OUTPUT_DIR)
```

```
print("Existe TEST_JSONL:", TEST_JSONL.exists())
```

```
print("Existe OUTPUT_DIR:", OUTPUT_DIR.exists())
```

```
TEST_JSONL: /content/drive/MyDrive/TT2_colab/data/sft_ready/feina_repr30_test_sft.jsonl
```

```
OUTPUT_DIR: /content/drive/MyDrive/TT2_colab/outputs/lora_runs
```

```
Existe TEST_JSONL: True
```

```
Existe OUTPUT_DIR: True
```

```
MODEL_EVALS = [
```

```
{
```

```
    "model_key": "llama3",
```

```
    "model_id": "meta-llama/Meta-Llama-3-8B-Instruct",
```

```
    "run_name": "lora_llama3_feina_repr30_v1",
```

```
},
```

```
{
```

```
    "model_key": "mistral",
```

```
    "model_id": "mistralai/Mistral-7B-Instruct-v0.2",
```

```
    "run_name": "lora_mistral_feina_repr30_v1",
```

```
},
```

```
]
```

```
display(pd.DataFrame(MODEL_EVALS))
```

	model_key	model_id	run_name	
0	llama3	meta-llama/Meta-Llama-3-8B-Instruct	lora_llama3_feina_repr30_v1	
1	mistral	mistralai/Mistral-7B-Instruct-v0.2	lora_mistral_feina_repr30_v1	

```
data_files = {
    "test": str(TEST_JSONL),
}
```

```
dataset = load_dataset("json", data_files=data_files)
```

```
test_ds = dataset["test"]
```

```
print("test_ds:", test_ds)
```

```
print(test_ds[0].keys())
```

```
print(test_ds[0]["instruction"][:500])
```

```
assert test_ds.num_rows == 238, "Test no coincide con repr30"
print("OK test repr30 cargado correctamente")
```

Generating test split: 238/0 [00:00<00:00, 6517.36 examples/s]

```
test_ds: Dataset({
  features: ['row_id', 'instruction', 'output', 'text'],
  num_rows: 238
})
```

```
dict_keys(['row_id', 'instruction', 'output', 'text'])
```

Reescribe en español el siguiente texto con lenguaje más claro y sencillo.
Conserva el significado original y no inventes información.

Devuelve solo la versión final simplificada.

Texto:

Muchas personas, más de las que uno pueda imaginarse, no tienen el hábito de presupue

Versión simplificada:

OK test repr30 cargado correctamente

```
def recover_source_text_from_instruction(instruction: str) -> str:
    text = str(instruction)
    marker = "Texto:"
    end_marker = "Versión simplificada:"

    if marker in text:
        text = text.split(marker, 1)[1]
    if end_marker in text:
        text = text.split(end_marker, 1)[0]

    text = text.replace("\n", " ").replace("\r", " ").replace("\t", " ")
    text = " ".join(text.split()).strip()
    return text
```

```
compute_dtype = torch.float16
```

```
bnb_config = BitsAndBytesConfig(
    load_in_4bit=True,
    bnb_4bit_quant_type="nf4",
    bnb_4bit_use_double_quant=True,
    bnb_4bit_compute_dtype=torch.float16,
)
```

```
bnb_config
```

```
BitsAndBytesConfig {
  "_load_in_4bit": true,
  "_load_in_8bit": false,
  "bnb_4bit_compute_dtype": "float16",
  "bnb_4bit_quant_storage": "uint8",
  "bnb_4bit_quant_type": "nf4",
  "bnb_4bit_use_double_quant": true,
  "llm_int8_enable_fp32_cpu_offload": false,
  "llm_int8_has_fp16_weight": false,
  "llm_int8_skip_modules": null,
  "llm_int8_threshold": 6.0,
  "load_in_4bit": true,
  "load_in_8bit": false,
```

```
"quant_method": "bitsandbytes"
}
```

```
def clean_lora_output(decoded_text: str) -> str:
    text = str(decoded_text).strip()
    marker = "Versión simplificada:"

    if marker in text:
        text = text.split(marker, 1)[-1].strip()

    text = text.replace("\n", " ").replace("\r", " ").replace("\t", " ")
    text = " ".join(text.split()).strip()
    return text
```

```
def load_lora_model(model_id: str, adapter_dir: Path):
    tokenizer = AutoTokenizer.from_pretrained(
        model_id,
        use_fast=True,
        token=os.environ.get("HF_TOKEN"),
    )

    if tokenizer.pad_token is None:
        tokenizer.pad_token = tokenizer.eos_token
    tokenizer.padding_side = "right"

    print("pad_token:", tokenizer.pad_token)
    print("eos_token:", tokenizer.eos_token)

    t0 = time.perf_counter()
    base_model = AutoModelForCausalLM.from_pretrained(
        model_id,
        quantization_config=bnb_config,
        dtype=torch.float16,
        device_map="auto",
        token=os.environ.get("HF_TOKEN"),
    )

    model = PeftModel.from_pretrained(
        base_model,
        adapter_dir,
    )
    model.eval()

    t1 = time.perf_counter()
    print(f"Tiempo cargando modelo + adapter: {t1 - t0:.2f} s")

    return tokenizer, model
```

```
def generate_simplification(model, tokenizer, instruction: str, max_input_length: int):
    inputs = tokenizer(
        instruction,
        return_tensors="pt",
        truncation=True,
        max_length=max_input_length,
    ).to(model.device)

    with torch.no_grad():
```

```

        outputs = model.generate(
            **inputs,
            max_new_tokens=max_new_tokens,
            do_sample=False,
            pad_token_id=tokenizer.eos_token_id,
        )

    decoded = tokenizer.decode(outputs[0], skip_special_tokens=True)
    cleaned = clean_lora_output(decoded)
    return cleaned

```

```

from pathlib import Path
import sys

DRIVE_ROOT = Path("/content/drive/MyDrive/TT2_colab")

if str(DRIVE_ROOT) not in sys.path:
    sys.path.append(str(DRIVE_ROOT))

print("DRIVE_ROOT:", DRIVE_ROOT)
print("src existe:", (DRIVE_ROOT / "src").exists())
print("metrics existe:", (DRIVE_ROOT / "src" / "evaluation" / "metrics.py").exists())
print("sys.path ok:", str(DRIVE_ROOT) in sys.path)

```

```

DRIVE_ROOT: /content/drive/MyDrive/TT2_colab
src existe: True
metrics existe: True
sys.path ok: True

```

```

from src.evaluation.metrics import evaluate_dataframe, summarize_metrics
print("Import OK")

```

```

Import OK

```

```

def evaluate_one_lora_run(run_cfg: dict, test_ds):
    model_key = run_cfg["model_key"]
    model_id = run_cfg["model_id"]
    run_name = run_cfg["run_name"]

    run_dir = OUTPUT_DIR / run_name
    adapter_dir = run_dir / "final_adapter"
    eval_dir = run_dir / "evaluation"
    eval_dir.mkdir(parents=True, exist_ok=True)

    print("\n" + "=" * 90)
    print("Evaluando corrida")
    print("model_key :", model_key)
    print("model_id  :", model_id)
    print("run_name   :", run_name)
    print("adapter_dir:", adapter_dir)
    print("=" * 90)

    if not adapter_dir.exists():
        raise FileNotFoundError(f"No existe adapter_dir: {adapter_dir}")

    tokenizer, model = load_lora_model(model_id, adapter_dir)

```

```

# prueba rápida
for i in range(2):
    row = test_ds[i]
    pred = generate_simplification(model, tokenizer, row["instruction"], max_inp
    print("-" * 90)
    print("ROW ID:", row["row_id"])
    print("REFERENCE:", row["output"][:250])
    print("PREDICTION:", pred[:250])

records = []
t0 = time.perf_counter()

for i, row in enumerate(test_ds):
    pred = generate_simplification(
        model=model,
        tokenizer=tokenizer,
        instruction=row["instruction"],
        max_input_length=1024,
        max_new_tokens=256,
    )

    records.append({
        "row_id": row["row_id"],
        "instruction": row["instruction"],
        "generated_text": pred,
    })

    if (i + 1) % 25 == 0:
        print(f"Procesados: {i + 1}/{len(test_ds)}")

t1 = time.perf_counter()
generation_minutes = (t1 - t0) / 60
print(f"Tiempo total de generación en test: {generation_minutes:.2f} min")

lora_test_df = pd.DataFrame(records)
lora_test_df["source_text"] = lora_test_df["instruction"].apply(recover_source_t

test_df_from_sft = pd.DataFrame(test_ds)
test_df_from_sft = test_df_from_sft.rename(columns={"output": "reference_text"})

lora_test_df = lora_test_df.merge(
    test_df_from_sft[["row_id", "reference_text"]],
    on="row_id",
    how="left",
)

print("Shape predicciones:", lora_test_df.shape)
display(lora_test_df[["row_id", "source_text", "reference_text", "generated_text

pred_path = eval_dir / "lora_test_predictions.csv"
lora_test_df.to_csv(pred_path, index=False, encoding="utf-8-sig")
print("Predicciones guardadas en:", pred_path)

evaluated_lora_test_df = evaluate_dataframe(
    lora_test_df,
    source_col="source_text",
    pred_col="generated_text",
    ref_col="reference_text",

```

```

        compute_bertscore=True,
        compute_sbert=False,
    )

display(evaluated_lora_test_df.head(3))

lora_summary = summarize_metrics(
    evaluated_lora_test_df,
    group_cols=[],
)
display(lora_summary)

evaluated_path = eval_dir / "lora_test_evaluated.csv"
summary_path = eval_dir / "lora_test_summary.csv"

evaluated_lora_test_df.to_csv(evaluated_path, index=False, encoding="utf-8-sig")
lora_summary.to_csv(summary_path, index=False, encoding="utf-8-sig")

print("Evaluated guardado en:", evaluated_path)
print("Summary guardado en:", summary_path)

main_metric_cols = [
    c for c in [
        "sari",
        "bertscore_f1",
        "bleu",
        "rouge1_f",
        "rouge2_f",
        "rougeL_f",
        "fernandez_huerta_pred",
        "inflesz_pred",
        "compression_ratio_eval",
        "exact_copy",
    ]
    if c in lora_summary.columns
]

main_metrics = lora_summary[main_metric_cols].copy()
display(main_metrics)

result_row = {
    "run_name": run_name,
    "model_key": model_key,
    "model_id": model_id,
    "generation_minutes": generation_minutes,
}

for col in lora_summary.columns:
    result_row[col] = lora_summary.iloc[0][col]

# liberar memoria
del model
gc.collect()
if torch.cuda.is_available():
    torch.cuda.empty_cache()

return result_row

```



```
all_eval_results = []

for run_cfg in MODEL_EVALS:
    result_row = evaluate_one_lora_run(run_cfg, test_ds)
    all_eval_results.append(result_row)

eval_results_df = pd.DataFrame(all_eval_results)
display(eval_results_df)
```


https://colab.research.google.com/drive/1dbRysVwPDOUx5R58I3hnvt98I9NvmzGT?authuser=2#scrollTo=bhgoRT_cUrwh&printMo... 12/24

Procesados: 100/238

Procesados: 125/238

search.google.com/drive/1dbBysVwPDQI1x5B5813bnvt98I9NvmzGT?authuser=2#scrollTo=bhgqBT_cI1rwb&printMo 13

[illegible]

Procesados: 225/238

Tiempo total de generación en test: 19.40 min

Shape predicciones: (238, 5)

config.json: 100%

625/625 [00:00<00:00, 64.6kB/s]

tokenizer_config.json: 100%

49.0/49.0 [00:00<00:00, 5.12kB/s]

vocab.txt: 996k/? [00:00<00:00, 1.34MB/s]

tokenizer.json: 1.96M/? [00:00<00:00, 2.14MB/s]

model.safetensors: 100%

714M/714M [00:05<00:00, 177MB/s]

Loading weights: 100%

199/199 [00:00<00:00, 3702.13it/s]

```
[transformers] BertModel LOAD REPORT from: bert-base-multilingual-cased
```

Key	Status		
cls.predictions.transform.dense.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.bias	UNEXPECTED		
cls.seq_relationship.weight	UNEXPECTED		
cls.predictions.transform.dense.bias	UNEXPECTED		
cls.predictions.bias	UNEXPECTED		
cls.seq_relationship.bias	UNEXPECTED		

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 3686.10it/s]

[transformers] **BertModel LOAD REPORT** from: bert-base-multilingual-cased

Key	Status		
cls.predictions.transform.dense.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.bias	UNEXPECTED		
cls.seq_relationship.weight	UNEXPECTED		
cls.predictions.transform.dense.bias	UNEXPECTED		
cls.predictions.bias	UNEXPECTED		
cls.seq_relationship.bias	UNEXPECTED		

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 3607.29it/s]

[transformers] **BertModel LOAD REPORT** from: bert-base-multilingual-cased

Key	Status		
cls.predictions.transform.dense.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.bias	UNEXPECTED		
cls.seq_relationship.weight	UNEXPECTED		
cls.predictions.transform.dense.bias	UNEXPECTED		
cls.predictions.bias	UNEXPECTED		
cls.seq_relationship.bias	UNEXPECTED		

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 3405.55it/s]

[transformers] **BertModel LOAD REPORT** from: bert-base-multilingual-cased

Key	Status		
cls.predictions.transform.dense.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.bias	UNEXPECTED		
cls.seq_relationship.weight	UNEXPECTED		
cls.predictions.transform.dense.bias	UNEXPECTED		
cls.predictions.bias	UNEXPECTED		
cls.seq_relationship.bias	UNEXPECTED		

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 3660.69it/s]

[transformers] **BertModel LOAD REPORT** from: bert-base-multilingual-cased

Key	Status		
cls.predictions.transform.dense.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.bias	UNEXPECTED		
cls.seq_relationship.weight	UNEXPECTED		
cls.predictions.transform.dense.bias	UNEXPECTED		
cls.predictions.bias	UNEXPECTED		
cls.seq_relationship.bias	UNEXPECTED		

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 3503.59it/s]

[transformers] **BertModel LOAD REPORT** from: bert-base-multilingual-cased

Key	Status		
cls.predictions.transform.dense.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.bias	UNEXPECTED		
cls.seq_relationship.weight	UNEXPECTED		
cls.predictions.transform.dense.bias	UNEXPECTED		
cls.predictions.bias	UNEXPECTED		
cls.seq_relationship.bias	UNEXPECTED		

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 3643.64it/s]

[transformers] BertModel LOAD REPORT from: bert-base-multilingual-cased

Key	Status		
cls.predictions.transform.dense.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.bias	UNEXPECTED		
cls.seq_relationship.weight	UNEXPECTED		
cls.predictions.transform.dense.bias	UNEXPECTED		
cls.predictions.bias	UNEXPECTED		
cls.seq_relationship.bias	UNEXPECTED		

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 4630.86it/s]

[transformers] BertModel LOAD REPORT from: bert-base-multilingual-cased

Key	Status		
cls.predictions.transform.dense.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.bias	UNEXPECTED		
cls.seq_relationship.weight	UNEXPECTED		
cls.predictions.transform.dense.bias	UNEXPECTED		
cls.predictions.bias	UNEXPECTED		
cls.seq_relationship.bias	UNEXPECTED		

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 3874.74it/s]

[transformers] BertModel LOAD REPORT from: bert-base-multilingual-cased

Key	Status		
cls.predictions.transform.dense.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.bias	UNEXPECTED		
cls.seq_relationship.weight	UNEXPECTED		
cls.predictions.transform.dense.bias	UNEXPECTED		
cls.predictions.bias	UNEXPECTED		
cls.seq_relationship.bias	UNEXPECTED		

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 6144.71it/s]

[transformers] BertModel LOAD REPORT from: bert-base-multilingual-cased

Key	Status		
cls.predictions.transform.dense.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.bias	UNEXPECTED		

cls.seq_relationship.weight	UNEXPECTED
cls.predictions.transform.dense.bias	UNEXPECTED
cls.predictions.bias	UNEXPECTED
cls.seq_relationship.bias	UNEXPECTED

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 3172.52it/s]

[transformers] BertModel LOAD REPORT from: bert-base-multilingual-cased

Key	Status
cls.predictions.transform.dense.weight	UNEXPECTED
cls.predictions.transform.LayerNorm.weight	UNEXPECTED
cls.predictions.transform.LayerNorm.bias	UNEXPECTED
cls.seq_relationship.weight	UNEXPECTED
cls.predictions.transform.dense.bias	UNEXPECTED
cls.predictions.bias	UNEXPECTED
cls.seq_relationship.bias	UNEXPECTED

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 3414.23it/s]

[transformers] BertModel LOAD REPORT from: bert-base-multilingual-cased

Key	Status
cls.predictions.transform.dense.weight	UNEXPECTED
cls.predictions.transform.LayerNorm.weight	UNEXPECTED
cls.predictions.transform.LayerNorm.bias	UNEXPECTED
cls.seq_relationship.weight	UNEXPECTED
cls.predictions.transform.dense.bias	UNEXPECTED
cls.predictions.bias	UNEXPECTED
cls.seq_relationship.bias	UNEXPECTED

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 4219.13it/s]

[transformers] BertModel LOAD REPORT from: bert-base-multilingual-cased

Key	Status
cls.predictions.transform.dense.weight	UNEXPECTED
cls.predictions.transform.LayerNorm.weight	UNEXPECTED
cls.predictions.transform.LayerNorm.bias	UNEXPECTED
cls.seq_relationship.weight	UNEXPECTED
cls.predictions.transform.dense.bias	UNEXPECTED
cls.predictions.bias	UNEXPECTED
cls.seq_relationship.bias	UNEXPECTED

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 4987.43it/s]

[transformers] BertModel LOAD REPORT from: bert-base-multilingual-cased

Key	Status
cls.predictions.transform.dense.weight	UNEXPECTED
cls.predictions.transform.LayerNorm.weight	UNEXPECTED
cls.predictions.transform.LayerNorm.bias	UNEXPECTED
cls.seq_relationship.weight	UNEXPECTED
cls.predictions.transform.dense.bias	UNEXPECTED
cls.predictions.bias	UNEXPECTED
cls.seq_relationship.bias	UNEXPECTED

```
notes:
- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i
Loading weights: 100% 199/199 [00:00<00:00, 3561.44it/s]
[transformers] BertModel LOAD REPORT from: bert-base-multilingual-cased
Key | Status | |
-----+-----+---+
cls.predictions.transform.dense.weight | UNEXPECTED | |
cls.predictions.transform.LayerNorm.weight | UNEXPECTED | |
cls.predictions.transform.LayerNorm.bias | UNEXPECTED | |
cls.seq_relationship.weight | UNEXPECTED | |
cls.predictions.transform.dense.bias | UNEXPECTED | |
cls.predictions.bias | UNEXPECTED | |
cls.seq_relationship.bias | UNEXPECTED | |
```

```
Notes:
- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i
```

row_id	instruction	generated_text	source_text	reference_text	src_words	pred_w
0	92	Reescribe en español el siguiente texto con le...	Muchas personas no tienen el hábito de presupu...	Muchas personas, más de las que uno pueda imag...	Muchas personas no tienen el hábito de presupu...	35
1	115	Reescribe en español el siguiente texto con le...	Uno de los problemas que enfrentan las persona...	Uno de los mayores problemas que enfrentan las...	Uno de los mayores problemas de las personas, ...	50
2	118	Reescribe en español el siguiente texto con le...	El comportamiento del ingreso marca etapas o c...	El comportamiento del ingreso marca etapas o c...	El comportamiento del ingreso marca etapas o c...	49

3 rows × 30 columns

sari	bleu	fernandez_huerta_pred	fernandez_huerta_src	fernandez_huerta_de
0	40.661328	45.025877	71.394097	76.959595

Evaluated guardado en: /content/drive/MyDrive/TT2_colab/outputs/lora_runs/lora_llama3
Summary guardado en: /content/drive/MyDrive/TT2_colab/outputs/lora_runs/lora_llama3_f

sari	bertscore_f1	bleu	rouge1_f	rouge2_f	rougeL_f	fernandez_huerta_pre
0	40.661328	0.90034	45.025877	0.734021	0.58785	0.691667

```
=====
Evaluando corrida
model_key : mistral
model_id : mistralai/Mistral-7B-Instruct-v0.2
run_name : lora_mistral_feina_repr30_v1
adapter_dir: /content/drive/MyDrive/TT2_colab/outputs/lora_runs/lora_mistral_feina_re
=====

config.json: 100% 596/596 [00:00<00:00, 40.0kB/s]

tokenizer_config.json: 2.10k/? [00:00<00:00, 130kB/s]

tokenizer.json: 1.80M/? [00:00<00:00, 51.5MB/s]

tokenizer.model: 100% 493k/493k [00:00<00:00, 2.48MB/s]

special_tokens_map.json: 100% 414/414 [00:00<00:00, 36.8kB/s]

pad_token: </s>
eos_token: </s>
```

model.safetensors.index.json: 25.1k/? [00:00<00:00, 1.02MB/s]

Download complete: 100% 14.5G/14.5G [09:55<00:00, 32.3MB/s]

Fetching 3 files: 100% 3/3 [09:55<00:00, 249.48s/it]

Loading weights: 100% 291/291 [01:05<00:00, 5.69it/s]

generation_config.json: 100% 111/111 [00:00<00:00, 10.9kB/s]

Tiempo cargando modelo + adapter: 669.04 s

ROW ID: 92

REFERENCE: Muchas personas no tienen el hábito de presupuestar sus gastos con base en

PREDICTION: Muchas personas no tienen el hábito de presupuestar sus gastos con base e

ROW ID: 115

REFERENCE: Uno de los mayores problemas de las personas, corporaciones, gobiernos y p

PREDICTION: Uno de los mayores problemas que enfrentan las personas, las corporacione

Procesados: 25/238

Procesados: 50/238

Procesados: 75/238

Procesados: 100/238

Procesados: 125/238

Procesados: 150/238

Procesados: 175/238

Procesados: 200/238

Procesados: 225/238

Tiempo total de generación en test: 22.90 min

Shape predicciones: (238, 5)

row_id		source_text	reference_text	generated_text
0	92	Muchas personas, más de las que uno pueda imag...	Muchas personas no tienen el hábito de presupu...	Muchas personas no tienen el hábito de presupu...
1	115	Uno de los mayores problemas que enfrentan las...	Uno de los mayores problemas de las personas, ...	Uno de los mayores problemas que enfrentan las...
2	118	El comportamiento del ingreso marca etapas o c...	El comportamiento del ingreso marca etapas o c...	La primera etapa es la de la educación básica,...

Predicciones guardadas en: /content/drive/MyDrive/TT2_colab/outputs/lora_runs/lora_mi

Loading weights: 100% 199/199 [00:00<00:00, 3153.76it/s]

[transformers] BertModel LOAD REPORT from: bert-base-multilingual-cased

Key	Status		
-----+-----+-----			
cls.predictions.transform.dense.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.bias	UNEXPECTED		
cls.seq_relationship.weight	UNEXPECTED		
cls.predictions.transform.dense.bias	UNEXPECTED		
cls.predictions.bias	UNEXPECTED		
cls.seq_relationship.bias	UNEXPECTED		

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100% 199/199 [00:00<00:00, 7352.53it/s]

[transformers] BertModel LOAD REPORT from: bert-base-multilingual-cased

Key	Status		
-----+-----+-----			
cls.predictions.transform.dense.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.bias	UNEXPECTED		
cls.seq_relationship.weight	UNEXPECTED		

cls.predictions.transform.dense.bias	UNEXPECTED	
cls.predictions.bias	UNEXPECTED	
cls.seq_relationship.bias	UNEXPECTED	

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 3452.36it/s]

[transformers] BertModel LOAD REPORT from: bert-base-multilingual-cased

Key	Status	
-----+-----+		
cls.predictions.transform.dense.weight	UNEXPECTED	
cls.predictions.transform.LayerNorm.weight	UNEXPECTED	
cls.predictions.transform.LayerNorm.bias	UNEXPECTED	
cls.seq_relationship.weight	UNEXPECTED	
cls.predictions.transform.dense.bias	UNEXPECTED	
cls.predictions.bias	UNEXPECTED	
cls.seq_relationship.bias	UNEXPECTED	

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 4101.78it/s]

[transformers] BertModel LOAD REPORT from: bert-base-multilingual-cased

Key	Status	
-----+-----+		
cls.predictions.transform.dense.weight	UNEXPECTED	
cls.predictions.transform.LayerNorm.weight	UNEXPECTED	
cls.predictions.transform.LayerNorm.bias	UNEXPECTED	
cls.seq_relationship.weight	UNEXPECTED	
cls.predictions.transform.dense.bias	UNEXPECTED	
cls.predictions.bias	UNEXPECTED	
cls.seq_relationship.bias	UNEXPECTED	

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 3551.17it/s]

[transformers] BertModel LOAD REPORT from: bert-base-multilingual-cased

Key	Status	
-----+-----+		
cls.predictions.transform.dense.weight	UNEXPECTED	
cls.predictions.transform.LayerNorm.weight	UNEXPECTED	
cls.predictions.transform.LayerNorm.bias	UNEXPECTED	
cls.seq_relationship.weight	UNEXPECTED	
cls.predictions.transform.dense.bias	UNEXPECTED	
cls.predictions.bias	UNEXPECTED	
cls.seq_relationship.bias	UNEXPECTED	

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 4258.09it/s]

[transformers] BertModel LOAD REPORT from: bert-base-multilingual-cased

Key	Status	
-----+-----+		
cls.predictions.transform.dense.weight	UNEXPECTED	
cls.predictions.transform.LayerNorm.weight	UNEXPECTED	
cls.predictions.transform.LayerNorm.bias	UNEXPECTED	
cls.seq_relationship.weight	UNEXPECTED	
cls.predictions.transform.dense.bias	UNEXPECTED	
cls.predictions.bias	UNEXPECTED	
cls.seq_relationship.bias	UNEXPECTED	

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 3006.02it/s]

[transformers] **BertModel LOAD REPORT** from: bert-base-multilingual-cased

Key	Status		
cls.predictions.transform.dense.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.bias	UNEXPECTED		
cls.seq_relationship.weight	UNEXPECTED		
cls.predictions.transform.dense.bias	UNEXPECTED		
cls.predictions.bias	UNEXPECTED		
cls.seq_relationship.bias	UNEXPECTED		

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 3911.42it/s]

[transformers] **BertModel LOAD REPORT** from: bert-base-multilingual-cased

Key	Status		
cls.predictions.transform.dense.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.bias	UNEXPECTED		
cls.seq_relationship.weight	UNEXPECTED		
cls.predictions.transform.dense.bias	UNEXPECTED		
cls.predictions.bias	UNEXPECTED		
cls.seq_relationship.bias	UNEXPECTED		

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 7594.78it/s]

[transformers] **BertModel LOAD REPORT** from: bert-base-multilingual-cased

Key	Status		
cls.predictions.transform.dense.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.bias	UNEXPECTED		
cls.seq_relationship.weight	UNEXPECTED		
cls.predictions.transform.dense.bias	UNEXPECTED		
cls.predictions.bias	UNEXPECTED		
cls.seq_relationship.bias	UNEXPECTED		

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 4085.33it/s]

[transformers] **BertModel LOAD REPORT** from: bert-base-multilingual-cased

Key	Status		
cls.predictions.transform.dense.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.weight	UNEXPECTED		
cls.predictions.transform.LayerNorm.bias	UNEXPECTED		
cls.seq_relationship.weight	UNEXPECTED		
cls.predictions.transform.dense.bias	UNEXPECTED		
cls.predictions.bias	UNEXPECTED		
cls.seq_relationship.bias	UNEXPECTED		

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 3160.77it/s]

[transformers] **BertModel LOAD REPORT** from: bert-base-multilingual-cased

Key	Status		
-----	--------	--	--

```

-----+-----+-----
cls.predictions.transform.dense.weight | UNEXPECTED | |
cls.predictions.transform.LayerNorm.weight | UNEXPECTED | |
cls.predictions.transform.LayerNorm.bias | UNEXPECTED | |
cls.seq_relationship.weight | UNEXPECTED | |
cls.predictions.transform.dense.bias | UNEXPECTED | |
cls.predictions.bias | UNEXPECTED | |
cls.seq_relationship.bias | UNEXPECTED | |

```

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 4178.68it/s]

[transformers] **BertModel LOAD REPORT** from: bert-base-multilingual-cased

```

Key | Status | |
-----+-----+-----
cls.predictions.transform.dense.weight | UNEXPECTED | |
cls.predictions.transform.LayerNorm.weight | UNEXPECTED | |
cls.predictions.transform.LayerNorm.bias | UNEXPECTED | |
cls.seq_relationship.weight | UNEXPECTED | |
cls.predictions.transform.dense.bias | UNEXPECTED | |
cls.predictions.bias | UNEXPECTED | |
cls.seq_relationship.bias | UNEXPECTED | |

```

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 4077.77it/s]

[transformers] **BertModel LOAD REPORT** from: bert-base-multilingual-cased

```

Key | Status | |
-----+-----+-----
cls.predictions.transform.dense.weight | UNEXPECTED | |
cls.predictions.transform.LayerNorm.weight | UNEXPECTED | |
cls.predictions.transform.LayerNorm.bias | UNEXPECTED | |
cls.seq_relationship.weight | UNEXPECTED | |
cls.predictions.transform.dense.bias | UNEXPECTED | |
cls.predictions.bias | UNEXPECTED | |
cls.seq_relationship.bias | UNEXPECTED | |

```

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 3487.60it/s]

[transformers] **BertModel LOAD REPORT** from: bert-base-multilingual-cased

```

Key | Status | |
-----+-----+-----
cls.predictions.transform.dense.weight | UNEXPECTED | |
cls.predictions.transform.LayerNorm.weight | UNEXPECTED | |
cls.predictions.transform.LayerNorm.bias | UNEXPECTED | |
cls.seq_relationship.weight | UNEXPECTED | |
cls.predictions.transform.dense.bias | UNEXPECTED | |
cls.predictions.bias | UNEXPECTED | |
cls.seq_relationship.bias | UNEXPECTED | |

```

Notes:

- UNEXPECTED: can be ignored when loading from different task/architecture; not ok i

Loading weights: 100%

199/199 [00:00<00:00, 4394.67it/s]

[transformers] **BertModel LOAD REPORT** from: bert-base-multilingual-cased

```

Key | Status | |
-----+-----+-----
cls.predictions.transform.dense.weight | UNEXPECTED | |
cls.predictions.transform.LayerNorm.weight | UNEXPECTED | |
cls.predictions.transform.LayerNorm.bias | UNEXPECTED | |
cls.seq_relationship.weight | UNEXPECTED | |

```



```
cls.predictions.transform.dense.bias
```

```
| UNEXPECTED |
```

```
global_eval_path = OUTPUT_DIR / "lora_eval_comparison_repr30.csv"  
eval_results_df.to_csv(global_eval_path, index=False, encoding="utf-8-sig")  
print("Comparación global guardada en:", global_eval_path)
```

```
- UNEXPECTED: can be ignored when loading from different task/architecture; not ok 1  
Comparación global guardada en: /content/drive/MyDrive/TT2_colab/outputs/lora_runs/lo  
row_id instruction generated_text source_text reference_text src_words pred_w
```

```
print("Contenido de OUTPUT_DIR:")  
for p in OUTPUT_DIR.iterdir():  
    print("-", p.name)
```